


Ξthereum - Yellow Paper

Outline

- Motivation
 - Blockchain
 - Block, State, Transactions
 - Gas, Payment
 - Contracts
 - Execution Model
 - Block Finalization
 - Applications
- 
- A decorative horizontal bar at the bottom of the slide, composed of several colored rectangular segments in shades of blue, green, teal, purple, orange, and yellow.

Motivation

*“What Ethereum intends to provide is a blockchain with a built-in fully fledged **Turing-complete programming language** which can be used to create **"contracts"** that encode **arbitrary state transition functions**, allowing users to create any of the systems described [below], as well as many others that we have not yet imagined, simply by writing up the logic in a few lines of code.”*

- *custom currencies and financial instruments (coloured coins)*
- *ownership of an underlying physical device (smart property)*
- *non-fungible assets (e.x. Namecoin)*
- *digital assets controlled by code (smart contracts)*
- *decentralized autonomous organizations (DAOs)*

- **Motivation**
- Blockchain
- Block, State, Transactions

Blockchain

- “Transaction-based state machine”
 - Account balances
 - Information of physical world
 - Etc.
- Transactions collected into blocks + previous block ID
- Miners dedicate effort (work) to bolster one block over a competitor's block
- Blocks include reward to incentivize miners

- Motivation
- **Blockchain**
- Block, State, Transactions

Blockchain

New state after Ethereum state transition function,

$$\sigma_{t+1} \equiv \Upsilon(\sigma_t, T)$$

Block as a collection of transactions,

$$B \equiv (... , (T_0 , T_1 , ...))$$

Block transition = state transition (transaction) + block finalization

$$\Pi(\sigma, B) \equiv \Omega(B, \Upsilon(\Upsilon(\sigma, T_0), T_1)...))$$

New state after block transition

$$\sigma_{t+1} \equiv \Pi(\sigma_t, B)$$

- Motivation
- **Blockchain**
- Block, State, Transactions


Blockchain

- **Decentralized** system; all miners have the opportunity to create a new block
- To form **consensus** over which **path** from genesis block to leaf to use an **agreed upon scheme** is required
- Disagreements = **forks**; kills **confidence** in the system
- **GHOST protocol** is the agreed-upon scheme

- Motivation
- **Blockchain**
- Block, State, Transactions

GHOST Protocol


- "Greedy Heaviest Observed Subtree"
- Fast confirmation times lead to **stale blocks** (wasted), and reduced security
- Exacerbated by miner centralization (pools)
- Include stale blocks ("Uncles") to determine longest ("heaviest") chain
- Uncle receives 87.5% of its base reward, nephew including it receives remaining 12.5%
- Only considers up to 7 generations



- Motivation
- **Blockchain**
- Block, State, Transactions

Block, State, Transactions: Account State

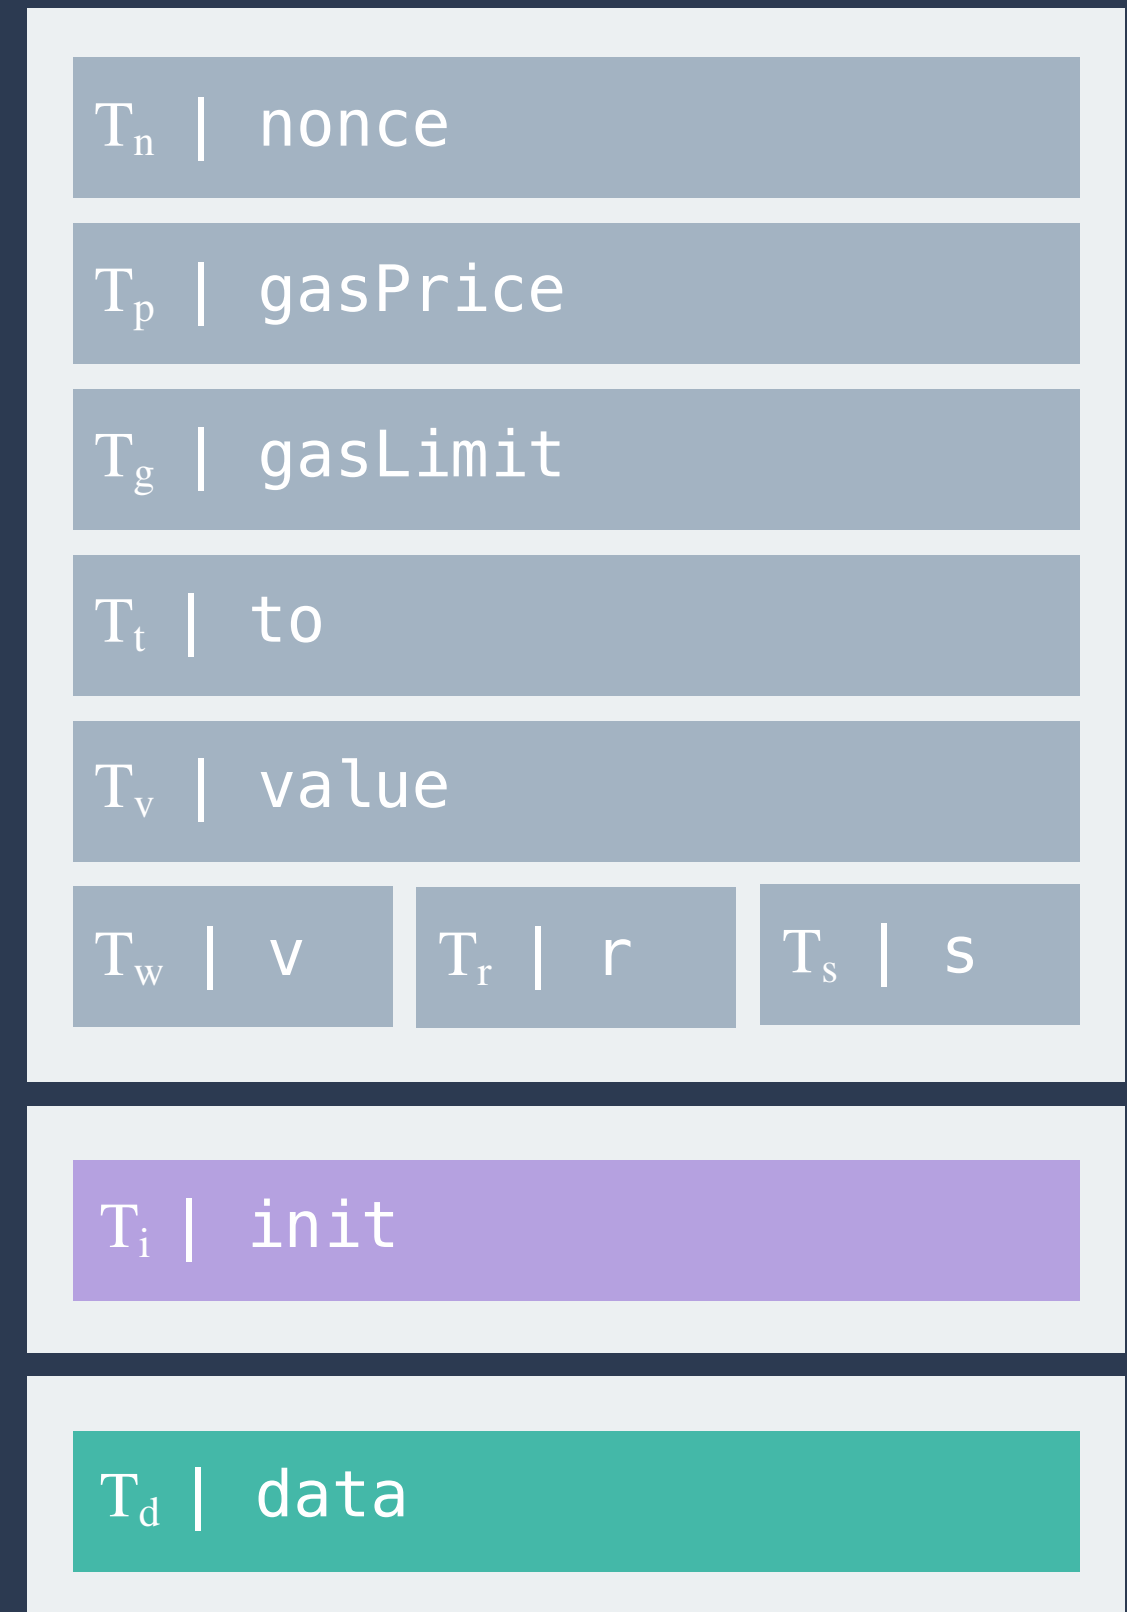
- **World state** (state tree) maps addresses (160-bit identifiers) and account states
- Mapping stored in a Merkle Patricia tree **off of the blockchain**
- Tree depends on database backend ("state database")



- Blockchain
- **Block, State, Transactions**
- Gas, Payment

Block, State, Transactions: The Transaction

- Two types of transactions
 1. **contract creation**
 2. **message calls**
- Resources cost gas: computation, and storage (transaction data)
- gasPrice, gasLimit put bounds on computation (to prevent DoS attacks)



Block, State, Transactions: The Block

- Block consists of header, **H**, transactions, **T**, ommer block headers, **U**

H _p parentHash	H _l gasLimit	
H _o ommerHash	H _g gasUsed	
H _c beneficiary	H _s timestamp	
H _r stateRoot	H _x extraData	
H _t transactionsRoot	H _m mixHash	H _n nonce
H _e receiptsRoot	T transactions ⋮	
H _b logsBloom		
H _d difficulty	U ommerBlockHeaders ⋮	
H _i number		

- Blockchain
- Block, State, Transactions**
- Gas, Payment

Block, State, Transactions: The Block

parentHash

- Hash of parent block's header

ommersHash

- Hash of ommer block header list

beneficiary

- Address where fees will be sent to

difficulty

- Derived from parent block's difficulty and timestamp

number

- Number of ancestor blocks

H _p parentHash	
H _o ommersHash	
H _c beneficiary	
H _d difficulty	
H _i number	

Block, State, Transactions: The Block

stateRoot

- Hash of root node of state tree after transactions are executed and finalizations applied
- $\sigma_{t+1} \equiv \Pi(\sigma_t, B)$

transactionsRoot

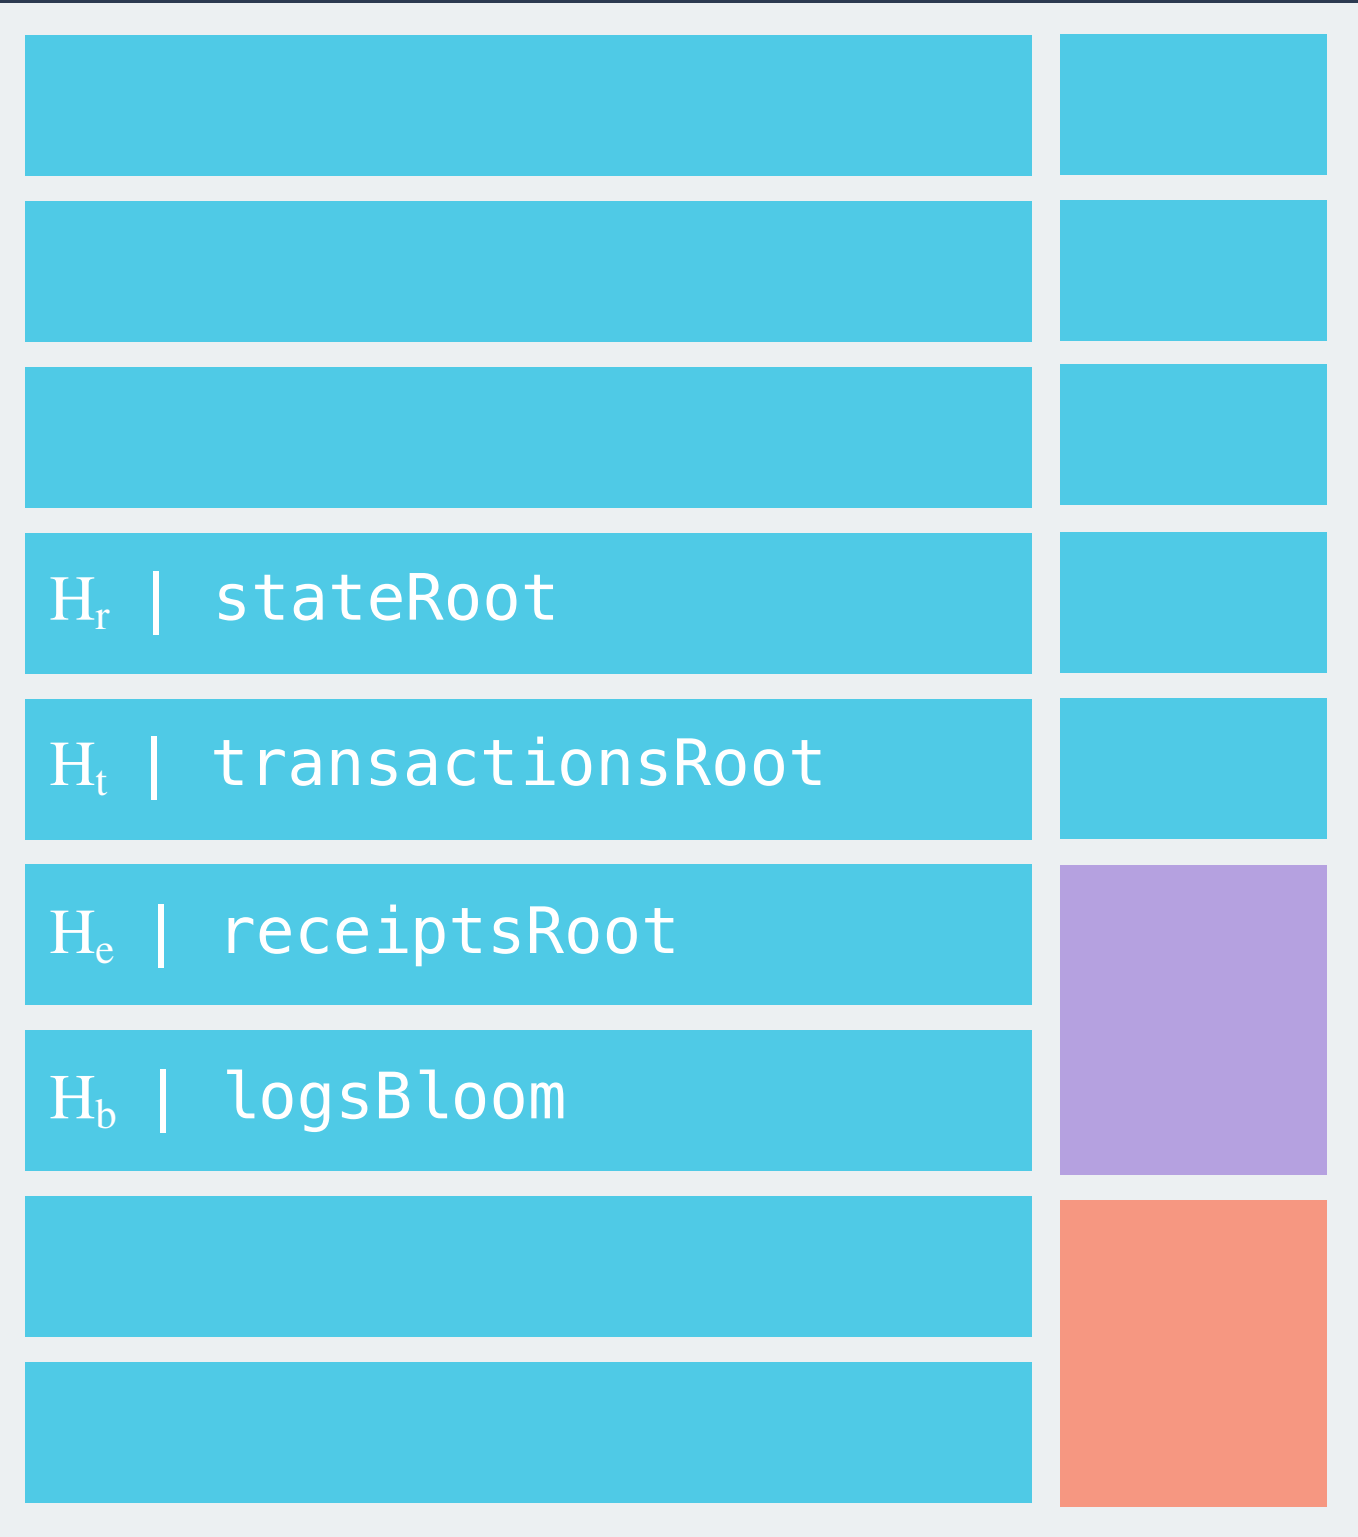
- Hash of root node of transaction tree

receiptsRoot

- Hash of root node of receipts tree populated from receipts of each transaction
- $L_R(R) \equiv (\text{TRIE}(L_S(R_\sigma)), R_u, R_b, R_l)$

logsBloom

- Bloom filter composed of indexable data in each log entry
- Node can quickly scan block headers check bloom filter if block may contain logs
- Re-executes transactions to generate logs, then returns



Block, State, Transactions: The Block

gasLimit

- Max gas all transactions combined are allowed to consume
- Akin to Bitcoin's blocksize (in bytes) but controllable by miner within constraints

gasUsed

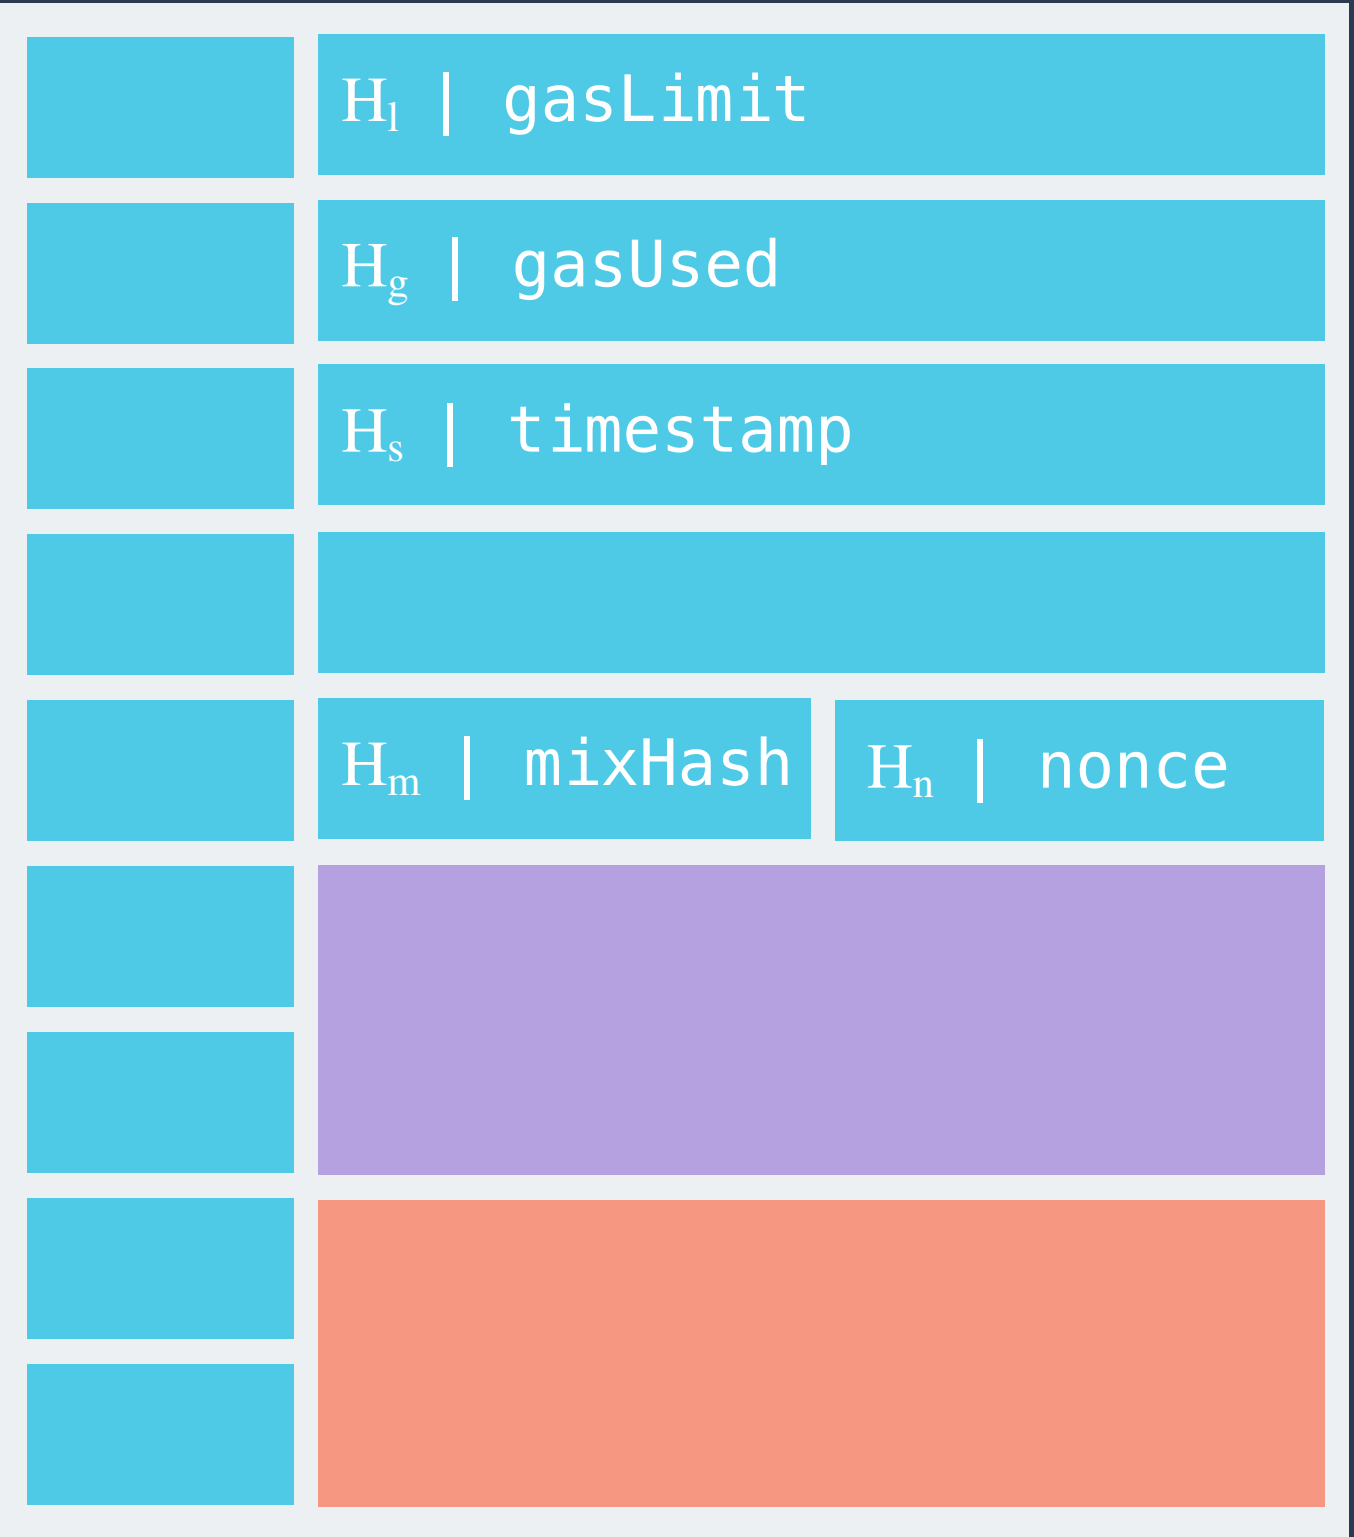
- Total gas used from executing transactions

timestamp

- Unix time() at this block's inception

mixHash, nonce

- Proof a sufficient amount of work went into mining this block



Gas, Payment

- Ether is used to purchase gas
- Transactors are free to set any **gasPrice** and miners are free to ignore any transactions
- **Intrinsic gas**, g_0 ("cost to pay for existence of transaction")
 - Amount of gas transaction requires to be paid prior to execution
 - Certain cost per operation (e.x. `SHA3(...)` costs 30 gas)
 - Remaining gas used to pay for computation $g \equiv T_g - g_0$
- Upfront transaction cost: **gasLimit * gasPrice + value**
- **Unused gas** ($*g$) is refunded
- In-case of **insufficient funds** during execution, state is reverted, but fees are kept by miner

- Block, State, Transactions
- **Gas, Payment**
- Contracts

Execution Model

- Specifies how system state is altered given
 - Bytecode instructions
 - Environment tuple data
- Ethereum Virtual Machine (EVM) - *quasi*-Turing-complete machine
- Stack-based, 256-bit word item size, 1024-item max stack size
- Word-addressable byte array memory model
- Word-addressable word array storage model (non-volatile, stored separately in system state, $K/V, \sigma$)
- Fees associated with memory and storage access

- Contracts
- **Execution Model**
- Block Finalization

Execution Model

Execution model defines,

$$(\sigma', g', A, o) \equiv \Xi(\sigma, g, I)$$

Which can compute the resultant state, σ' , remaining gas, g' , accrued substate, A , and resultant output, o

Execution Environment

σ | System state

g | Remaining gas for computation

I_a | Code-owning account address

I_o | Transactor (sender) address

I_p | Gas price of transaction

I_d | Input data byte array

I_s | Execution-initiator address (E.x. I_o)

I_v | Value in Wei sent to account

I_b | Machine code to be executed

I_H | Block header of present block

I_e | Depth of present message-call

Execution Model

- Stack items are added/removed
- Single run loop,
 - Check *exception*-halting state at each step (Eq. 128)
 - Program counter increments for each operation
 - Gas is reduced by the instruction's gas cost
 - Check *normal*-halting state at each step (Eq. 132)
- JUMPDEST command allows jumping to arbitrary positions in contract code
- Appendix H.2 for complete instruction set

Block Finalization

Finalizing a block involves,

1. Validate (miners: determine) omomers
 - Valid headers, ≤ 7 generations
2. Validate (miners: determine) transactions
 - Gas used in block (gasUsed) == accumulated gas used according to final transaction
3. Apply rewards
 - Reward miner 5 ether + $1/32 * 5$ ether per ommer block
 - Reward ommer depending on block number
4. Verify (miners: compute valid) state and nonce

- Execution Model
- **Block Finalization**
- Applications

Block Finalization - Mining

- The mining PoW exists as a cryptographically-secure nonce proving a particular amount of computation was expended
- Mining new blocks comes with an attached reward to incentivize work, and distribute wealth
- Goals of a PoW function
 - Accessible, not requiring specialized hardware
 - Prevent super-linear profits, and gaining large % of mining power
 - ASIC-resistant

- Execution Model
- **Block Finalization**
- Applications

Block Finalization - Ethash

- Ethash is the PoW algorithm for Ethereum
- Algorithm,
 1. Compute a seed by scanning through block headers up until current block
 2. Compute 16MB pseudorandom cache (stored by light clients)
 3. Generate 1 GB dataset (DAG) from cache, where dataset item depends on small number of cache items - updated every 30k blocks ~ 100Hrs - grows linearly with time
 4. Mining involves hashing together random slices of the dataset - memory hard task

- Execution Model
- **Block Finalization**
- Applications

Applications

- Financial applications
 - sub-currencies, financial derivatives, hedging contracts, wallets, wills, employment contracts, insurance
- Online voting, gambling, decentralized governance
- Decentralized storage, cloud computing
- “Decentralized Autonomous Organizations”
- Smart property
- Dapps (e.x. Golem, Augur)
- ICOs

- Execution Model
- Block Finalization
- **Applications**

Applications - Developer Tools

- **Writing smart contracts (languages)**

- Solidity
- Serpent
- LLL (Low-level Lisp-like Language)

- **Clients**

- Geth
- testRPC

- **Frameworks**

- Truffle
- OpenZeppelin

- **Dapps**

- Web3.js
- MetaMask
- Mist

- Execution Model
- Block Finalization
- **Applications**

Solidity Example - MetaCoin.sol

```
pragma solidity ^0.4.4;

import "../ConvertLib.sol";

contract MetaCoin {
    mapping (address => uint) balances;

    event Transfer(address indexed _from, address indexed _to, uint256 _value);

    function MetaCoin() {
        balances[tx.origin] = 10000;
    }

    function sendCoin(address receiver, uint amount) returns(bool sufficient) {
        if (balances[msg.sender] < amount) return false;
        balances[msg.sender] -= amount;
        balances[receiver] += amount;
        Transfer(msg.sender, receiver, amount);
        return true;
    }

    function getBalanceInEth(address addr) returns(uint){
        return ConvertLib.convert(getBalance(addr), 2);
    }

    function getBalance(address addr) returns(uint) {
        return balances[addr];
    }
}
```